

# DATA MENTORSHIP PROGRAM

*This document teaches basic and advanced statistical problem-solving using python. Please ensure that you have downloaded Python on your computer. The datasets are saved in Google sheet, please download the sheet on your computer before uploading the dataset to Python*

Owner: `mdave5@horizon.csueastbay.edu` Rajvi Mehta

## Basic Statistics using Python

1. Statistics in Python - why is it important?
2. Introduction to NumPy and Pandas and other common python libraries for basic statistics
3. Basic Descriptive analysis in Python (mean, median, mode, standard deviation)
4. Advanced Statistics in Python (Hypothesis testing, introduction to pvalue, probability, and sampling distributions)
5. Simple Linear regression

## Statistics in Python

Statistics refers to the *art and science* of collecting, analyzing, presenting and interpreting data. As a data professional, your knowledge of statistics will be crucial as you progress through your data career. Statistical knowledge will help you understand the outcome of your analysis/research better while helping to solve your client's problem.

This exercise will help you conduct basic and advanced statistical processes using Python. To get a better understanding of some of the most commonly used statistical terms, refer to this [blog](#).

## Choosing your Python Statistics Libraries

1. **Python Statistics:** Built-in library for descriptive statistics. Use it when your dataset is not too large.
2. **Numpy:** Third-party library for numerical computing.
3. **Scipy:** This library is based on Numpy. It offers additional functionality (including scipy.stats) for statistical analysis.
4. **Pandas:** This is a third-party library based on Numpy. Great for handling labeled one-dimensional data with series and two-dimensional data like DataFrame.
5. **Matplotlib:** Used mainly for data visualization.
6. **Seaborn:** Data visualization library built on top of Matplotlib to create statistical graphs. Used to explore data relationships and machine learning models.

## Part 3 - Exercise for Basic Statistics in Python

### Basic Descriptive analysis in Python (mean, median, mode, standard deviation)

**Owner:** mdave5@horizon.csueastbay.edu

**Goal:** This document shares a step-by-step guide on conducting descriptive statistics using Jupyter (Python script)

**Recommended to go through [this blog](#) for people who do not have prior knowledge of statistics.**

**Data Set:** We will be using the [Top 50 Billionaires dataset](#) for this project.

**Data Dictionary:** This dataset captures ranks, names, net worth, ages, financial sources, and countries of the top 50 billionaires.

- Rank - Numeric
- Name - Text
- Net Worth - Currency / Numeric
- Age - Numeric
- Financial Source - Text
- Countries - Text

### Data Questions to answer for this exercise

1. What are the average and median ages of the top 50 billionaires? (data. describe) (**basic EDA**)
2. What are the common trends among the top 10 billionaires and compare them to the bottom 10 billionaires? (bar graph, stacked bar chart)

3. Which continent have the most billionaires? (group by)
4. What is the correlation between age and net worth? Graphical Representation (scatter plot)
5. Which COUNTRY has the highest proportion of billionaires? Which country has the highest probability of billionaires whose net worth will be above average(group by sources)
6. ANOVA F-test (United States vs. other countries)

**1) Basic descriptive analysis (includes pictures)**

- **Reading data -**

```

: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
df=pd.read_csv('TOP 50 BILLIONAIRES LIST.csv')
df

```

|    | RANK | NAME                                  | NET WORTH | AGE  | SOURCE                      | COUNTRY/TERRITORY |
|----|------|---------------------------------------|-----------|------|-----------------------------|-------------------|
| 0  | 1.0  | Bernard Arnault & family              | \$176.6 B | 73.0 | LVMH                        | France            |
| 1  | 2.0  | Elon Musk                             | \$146.5 B | 51.0 | Tesla, SpaceX               | United States     |
| 2  | 3.0  | Gautam Adani                          | \$116.7 B | 60.0 | infrastructure, commodities | India             |
| 3  | 4.0  | Jeff Bezos                            | \$108.5 B | 58.0 | Amazon                      | United States     |
| 4  | 5.0  | Warren Buffett                        | \$106.3 B | 92.0 | Berkshire Hathaway          | United States     |
| 5  | 6.0  | Bill Gates                            | \$103.5 B | 67.0 | Microsoft                   | United States     |
| 6  | 7.0  | Larry Ellison                         | \$101.4 B | 78.0 | Oracle                      | United States     |
| 7  | 8.0  | Mukesh Ambani                         | \$87.9 B  | 65.0 | diversified                 | India             |
| 8  | 9.0  | Carlos Slim Helu & family             | \$83.6 B  | 82.0 | telecom                     | Mexico            |
| 9  | 10.0 | Steve Ballmer                         | \$78.2 B  | 66.0 | Microsoft                   | United States     |
| 10 | 11.0 | Larry Page                            | \$78.2 B  | 49.0 | Google                      | United States     |
| 11 | 12.0 | Michael Bloomberg                     | \$76.8 B  | 80.0 | Bloomberg LP                | United States     |
| 12 | 13.0 | Sergey Brin                           | \$75.0 B  | 49.0 | Google                      | United States     |
| 13 | 14.0 | Francoise Bettencourt Meyers & family | \$71.4 B  | 69.0 | L'Oréal                     | France            |
| 14 | 15.0 | Zhong Shanshan                        | \$67.8 B  | 68.0 | beverages, pharmaceuticals  | China             |
| 15 | 16.0 | Amancio Ortega                        | \$64.0 B  | 86.0 | Zara                        | Spain             |
| 16 | 17.0 | Jim Walton                            | \$61.5 B  | 74.0 | Walmart                     | United States     |
| 17 | 18.0 | Rob Walton                            | \$60.2 B  | 78.0 | Walmart                     | United States     |
| 18 | 19.0 | Alice Walton                          | \$59.3 B  | 73.0 | Walmart                     | United States     |
| 19 | 20.0 | Julia Koch & family                   | \$58.2 B  | 60.0 | Koch Industries             | United States     |
| 20 | 20.0 | Charles Koch                          | \$58.2 B  | 87.0 | Koch Industries             | United States     |

**Question 1:** What are the average and median ages of the top 50 billionaires?

- **Basic EDA**

```
: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 50 entries, 0 to 49
```

```
Data columns (total 6 columns):
```

| # | Column            | Non-Null Count | Dtype   |
|---|-------------------|----------------|---------|
| 0 | RANK              | 50 non-null    | float64 |
| 1 | NAME              | 50 non-null    | object  |
| 2 | NET WORTH         | 50 non-null    | object  |
| 3 | AGE               | 49 non-null    | float64 |
| 4 | SOURCE            | 50 non-null    | object  |
| 5 | COUNTRY/TERRITORY | 50 non-null    | object  |

```
dtypes: float64(2), object(4)
```

```
memory usage: 2.5+ KB
```

```
: df.describe()
```

```
:
```

|              | RANK     | AGE       |
|--------------|----------|-----------|
| <b>count</b> | 50.00000 | 49.000000 |
| <b>mean</b>  | 25.46000 | 68.816327 |
| <b>std</b>   | 14.56752 | 14.895012 |
| <b>min</b>   | 1.00000  | 38.000000 |
| <b>25%</b>   | 13.25000 | 58.000000 |
| <b>50%</b>   | 25.50000 | 69.000000 |
| <b>75%</b>   | 37.00000 | 82.000000 |
| <b>max</b>   | 50.00000 | 94.000000 |

The average (mean) and the median age of the top 50 billionaires are 25.46 years and 25.5 years respectively. This implies that the dataset is roughly symmetric and normally distributed as the mean is almost equal to the median.

**Question 2:** What are the common trends among the top 10 billionaires and compare them to the bottom 10 billionaires?

- Comparison by age

```
# Separate the top 10 and bottom 10 billionaires
bottom_10 = df.nlargest(10, "RANK")
top_10 = df.nsmallest(10, "RANK")

# Compare the mean age for each group
print("Top 10 Billionaires:")
print(top_10["AGE"].mean())
print()
print("Bottom 10 Billionaires:")
print(bottom_10["AGE"].mean())
print()

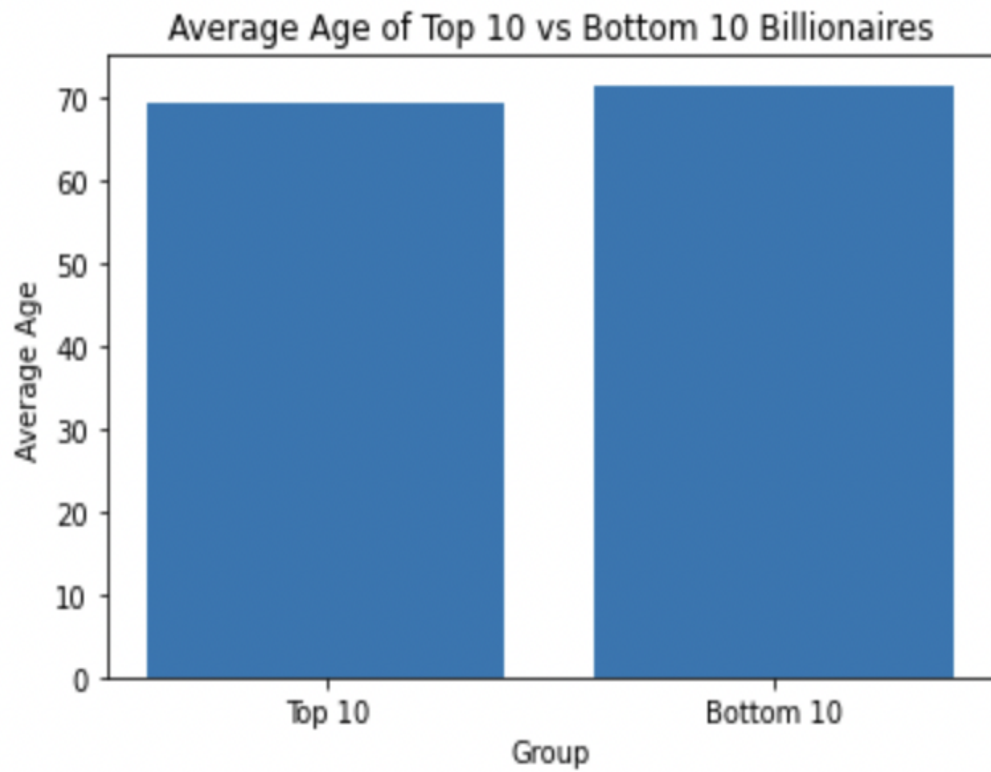
# Plot a bar chart of the average ages of the top 10 and bottom 10 billionaires
import matplotlib.pyplot as plt

plt.bar(["Top 10", "Bottom 10"], [top_10["AGE"].mean(), bottom_10["AGE"].mean()])
plt.xlabel("Group")
plt.ylabel("Average Age")
plt.title("Average Age of Top 10 vs Bottom 10 Billionaires")
plt.show()
```

---

Top 10 Billionaires:  
69.2

Bottom 10 Billionaires:  
71.4





- Comparison by country

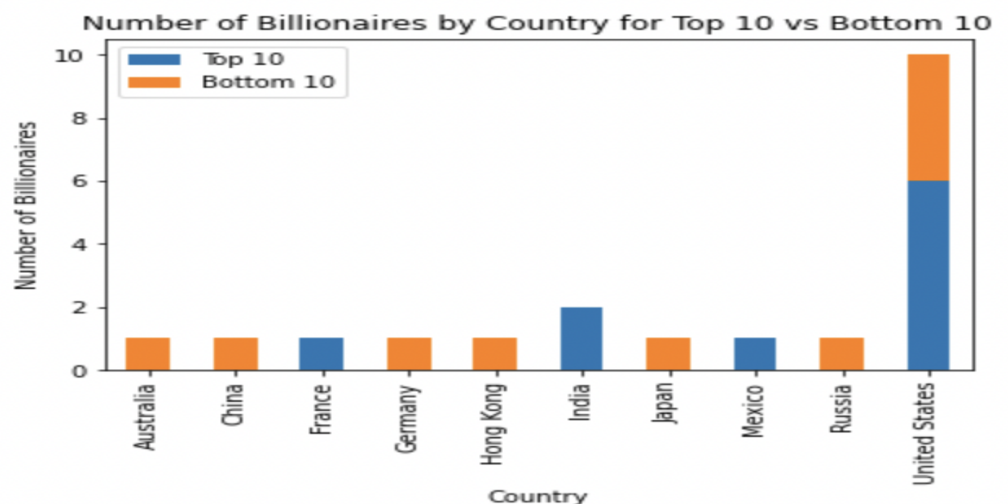
```
# Get the number of billionaires by country for each group
print("Top 10 Billionaires by Country:")
print(top_10["COUNTRY/TERRITORY"].value_counts())
print()
print("Bottom 10 Billionaires by Country:")
print(bottom_10["COUNTRY/TERRITORY"].value_counts())
top_10_counts = top_10["COUNTRY/TERRITORY"].value_counts()
bottom_10_counts = bottom_10["COUNTRY/TERRITORY"].value_counts()

# Combine the two series into a single DataFrame for plotting
counts = pd.DataFrame({"Top 10": top_10_counts, "Bottom 10": bottom_10_counts})
counts = counts.fillna(0).astype(int)

# Plot the stacked bar chart
counts.plot(kind="bar", stacked=True)
plt.xlabel("Country")
plt.ylabel("Number of Billionaires")
plt.title("Number of Billionaires by Country for Top 10 vs Bottom 10")
plt.legend(loc="upper left")
plt.show()
```

```
Top 10 Billionaires by Country:
United States      6
India              2
France            1
Mexico            1
Name: COUNTRY/TERRITORY, dtype: int64
```

```
Bottom 10 Billionaires by Country:
United States      4
Russia            1
Australia          1
Germany            1
China              1
Japan              1
Hong Kong          1
Name: COUNTRY/TERRITORY, dtype: int64
```





### Question 3: Which country has the most billionaires?

```
# Group the data by country/territory and count the number of billionaires in each group
grouped = df.groupby("COUNTRY/TERRITORY")["RANK"].count()

# Find the country/territory with the most billionaires
most_billionaires_country = grouped.idxmax()

print("The country/territory with the most billionaires is:", most_billionaires_country)
```

The country/territory with the most billionaires is: United States

### Question 4: What is the correlation between age and net worth?

Since, net worth is a combination of currency, numbers, and strings, we first need to remove the '\$' and 'B' characters from the entire Net Worth column to convert it to float to do any computations.

```
# replace the values and change the data type of the 'NET WORTH' column
df['NET WORTH'] = df['NET WORTH'].replace({'\$: ': '', 'B': ''}, regex=True).astype(float)

# Find correlation between 'AGE' and 'NET WORTH'
correlation = df['AGE'].corr(df['NET WORTH'])


print("The correlation between age and net worth of the top 50 billionaires is:", correlation)
```

The correlation between age and net worth of the top 50 billionaires is: -0.05763934928266901

A correlation coefficient of -0.058 indicates a small negative linear relationship between age and net worth. This suggests that there is a weak linear relationship between the two variables and that changes in one variable are not strongly associated with changes in the other. In this case, the negative sign indicates that as age increases, the net worth decreases, although the magnitude of the correlation is small.

```
import seaborn as sns

# Plot the scatterplot
plt.figure(figsize=(8, 6))
sns.scatterplot(x="AGE", y="NET WORTH", data=df)
plt.show()
```



**Question 5.** Which country has the highest proportion of billionaires?

```
: # Group the data by source and count the number of billionaires from each source
country_group = df.groupby('COUNTRY/TERRITORY').size().reset_index(name='counts')

# Calculate the proportion of billionaires from each source
country_group['proportion'] = country_group['counts'] / df.shape[0]

# Print the proportion of billionaires from each source
print(country_group[['COUNTRY/TERRITORY', 'proportion']])
```

|    | COUNTRY/TERRITORY | proportion |
|----|-------------------|------------|
| 0  | Australia         | 0.02       |
| 1  | Canada            | 0.02       |
| 2  | China             | 0.08       |
| 3  | France            | 0.12       |
| 4  | Germany           | 0.06       |
| 5  | Hong Kong         | 0.06       |
| 6  | India             | 0.04       |
| 7  | Italy             | 0.02       |
| 8  | Japan             | 0.02       |
| 9  | Mexico            | 0.02       |
| 10 | Russia            | 0.02       |
| 11 | Spain             | 0.02       |
| 12 | United States     | 0.50       |

Alternatively, to find which country has the highest probability of billionaires whose net worth will be above average (among the top 50), we can calculate the average net worth of all billionaires and then count the number of billionaires in each country whose net worth is above the average. Finally, divide this number by the total number of billionaires in each country to get the probability of having above-average billionaires in that country.

```
# Calculate the average net worth of all billionaires
average_net_worth = df['NET WORTH'].mean()

# Group the data by country
grouped = df.groupby(by='COUNTRY/TERRITORY')

# For each country, count the number of billionaires with net worth above the average
above_average_counts = grouped.apply(lambda x: x[x['NET WORTH'] > average_net_worth].count()['NET WORTH'])

# Calculate the total number of billionaires in each country
total_counts = grouped.count()['NET WORTH']

# Calculate the probability of having above average billionaires by dividing the number of above average billionaire
probabilities = above_average_counts / total_counts

# Find the country with the highest probability
result = probabilities.idxmax()
print("The country with the highest probability of having above average billionaires is:", result)
```

The country with the highest probability of having above average billionaires is: India

## Part 4 Advanced Statistical Analysis using Python

**Advanced Statistics in Python** (Hypothesis testing, introduction to p-value, probability, and sampling distributions)

**Owner:** **Rajvi Mehta**

**Goal:** This document shares a step-by-step guide on testing hypotheses in Jupyter (Python)

**Data Set:** We will be using the [50 Start-Ups dataset](#) for this project.

Please read [this blog](#) before proceeding.

**Data Dictionary:** This dataset captures R&D Spend, Administration cost, Marketing Spend, State, and Profit for 50 startups.

All data is in numeric, currencies.

Topics covered in this: EDA, Testing for Normality, T-test, Hypothesis testing

### Step 1: Upload the dataset

```
[7]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import linear_model
import statsmodels.api as sm
from sklearn.preprocessing import StandardScaler
%matplotlib inline
import statsmodels.api as sm
from datetime import datetime
import scipy.stats as stats
from sqlalchemy import Table, Column, Float, Integer, String, MetaData, ForeignKey
%matplotlib inline
sns.set()
```

```
[8]: #readingdataset
```

```
df = pd.read_csv('50_Startups.csv')
```

```
[9]: df.head()
```

|   | R&D Spend | Administration | Marketing Spend | State      | Profit    |
|---|-----------|----------------|-----------------|------------|-----------|
| 0 | 165349.20 | 136897.80      | 471784.10       | New York   | 192261.83 |
| 1 | 162597.70 | 151377.59      | 443898.53       | California | 191792.06 |
| 2 | 153441.51 | 101145.55      | 407934.54       | Florida    | 191050.39 |
| 3 | 144372.41 | 118671.85      | 383199.62       | New York   | 182901.99 |
| 4 | 142107.34 | 91391.77       | 366168.42       | Florida    | 166187.94 |

## Step 2: Conduct basic EDA

[5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
#   Column             Non-Null Count  Dtype  
---  -
0   R&D Spend          50 non-null     float64
1   Administration     50 non-null     float64
2   Marketing Spend    50 non-null     float64
3   State              50 non-null     object  
4   Profit             50 non-null     float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

[6]: `df.describe()`


|       | R&D Spend     | Administration | Marketing Spend | Profit        |
|-------|---------------|----------------|-----------------|---------------|
| count | 50.000000     | 50.000000      | 50.000000       | 50.000000     |
| mean  | 73721.615600  | 121344.639600  | 211025.097800   | 112012.639200 |
| std   | 45902.256482  | 28017.802755   | 122290.310726   | 40306.180338  |
| min   | 0.000000      | 51283.140000   | 0.000000        | 14681.400000  |
| 25%   | 39936.370000  | 103730.875000  | 129300.132500   | 90138.902500  |
| 50%   | 73051.080000  | 122699.795000  | 212716.240000   | 107978.190000 |
| 75%   | 101602.800000 | 144842.180000  | 299469.085000   | 139765.977500 |
| max   | 165349.200000 | 182645.560000  | 471784.100000   | 192261.830000 |

```
[41]: df.corr()
```

```
[41]:
```

|                 | R&D Spend | Administration | Marketing Spend | Profit   |
|-----------------|-----------|----------------|-----------------|----------|
| R&D Spend       | 1.000000  | 0.241955       | 0.724248        | 0.972900 |
| Administration  | 0.241955  | 1.000000       | -0.032154       | 0.200717 |
| Marketing Spend | 0.724248  | -0.032154      | 1.000000        | 0.747766 |
| Profit          | 0.972900  | 0.200717       | 0.747766        | 1.000000 |

```
[40]: sns.heatmap(df.corr());
```



### Step 3: Check Normality

What is Normality in statistics? The normal distribution describes the shape of the data (normally distributed data has a bell curve shape). Here, we are using a hypothesis test for checking the normality of the data.

Null Hypothesis: The data is not normally distributed.

Alternative Hypothesis: The data is normally distributed,

```
[8]: df2 = df[['R&D Spend', 'Administration', 'Marketing Spend', 'Profit']]

[9]: df2.columns

[9]: Index(['R&D Spend', 'Administration', 'Marketing Spend', 'Profit'], dtype='object')
```

## Check Normality

```
[11]: def check_normality(df2):
      test_stat_normality, p_value_normality = stats.shapiro(df2)
      print(p_value_normality)

[12]: group1 = df2[['R&D Spend']]
      group2 = df2[['Profit']]
      group3 = df2[['Marketing Spend']]

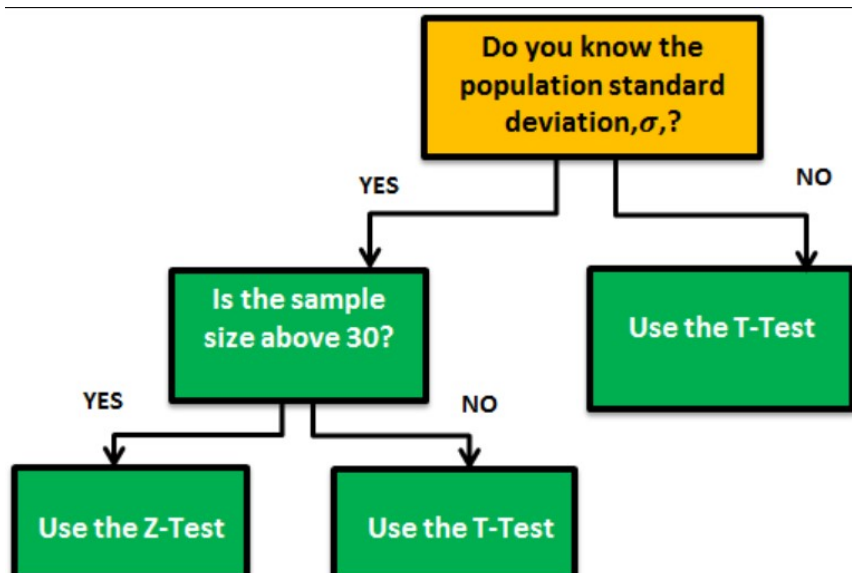
[13]: check_normality(group1)
      check_normality(group2)

0.18005183339118958
0.766564130783081
```

Fail to reject Null hypothesis >> the data is normally distributed

## Step 4: Conduct T Test

There are two types of tests that you can take at this point. T Test, Z test





## Conducting T Test since we have an unknown variance

```
[14]: ttest_p_value = stats.ttest_ind(group1,group2)

[15]: print(p_value)

[2.43343045e-05]
```

PValue is less than 0.05, hence we reject the null hypothesis - R&D spend drops profit

Step 5: Expand the Hypothesis to all three states.

## Testing the hypothesis between the three cities

```
[18]: df_ny = df[df['State']== 'New York']
      df_cl = df[df['State']== 'California']
      df_fl = df[df['State']== 'Florida']

[21]: ##Testing for New York

      dfny = df_ny[['R&D Spend','Administration','Marketing Spend','Profit']]

[32]: groupny1 = dfny[['R&D Spend']]
      groupny2 = dfny[['Profit']]

[33]: ttest_p_value = stats.ttest_ind(groupny1,groupny2)

[34]: print(p_value)

[0.02192342]

[35]: dfcl = df_cl[['R&D Spend','Administration','Marketing Spend','Profit']]
      groupcl1 = dfcl[['R&D Spend']]
      groupcl2 = dfcl[['Profit']]

[36]: ttest_p_value = stats.ttest_ind(groupcl1,groupcl2)
      print(p_value)

[0.01729946]

[37]: dffl = df_fl[['R&D Spend','Administration','Marketing Spend','Profit']]
      groupfl1 = dffl[['R&D Spend']]
      groupfl2 = dffl[['Profit']]

[38]: ttest_p_value = stats.ttest_ind(groupfl1,groupfl2)
      print(p_value)
```

```
[38]: ttest_p_value = stats.ttest_ind(groupfl1,groupfl2)
      print(p_value)

[0.01008887]
```

For all three states, the null hypothesis is false, which means that increase in R&D spend increases Profit